

Fei Protocol

Date	January 2021	
Lead Auditor	Valentin Wüstholz	
Co-auditors	Alexander Wade, Sergii Kravchenko	
Download	PDF 🚨	

1 Executive Summary

This report presents the results of our engagement with Fei Protocol to review some of the smart contracts in their stable coin implementation.

The review has been conducted over two weeks, from Jan 25, 2021 to Jan 29, 2021, by Valentin Wüstholz, Alexander Wade, and Sergii Kravchenko. Additionally, an infrastructure security assessment has been conducted over two weeks from Feb 8, 2021 to Feb 19, 2021, by Dominik Muhs. A total of 25 person-days were spent.

2 Scope

Our review focused on the commit hash ff892c5d. The list of files in scope and the priorities of the audit are defined by the client and can be found here.

The infrastructure assessment focused on the following assets:

- The fei.money domain, specifically ropsten-app.fei.finance
- The icaruscryptolab AWS organization
- The fei-protocol/fei-app at commit hash eee5d29

3 Findings

Each issue has an assigned severity:

- Minor issues are subjective in nature. They are typically suggestions around best practices or readability. Code maintainers should use their own judgment as to whether to address such issues.
- Medium issues are objective in nature but are not security vulnerabilities. These should be addressed unless there is a clear reason not to.

- Major issues are security vulnerabilities that may not be directly exploitable or may require certain conditions in order to be exploited. All major issues should be addressed.
- Critical issues are directly exploitable security vulnerabilities that need to be fixed.

3.1 GenesisGroup.commit overwrites previously-committed values (ritical)

Resolution

This was addressed in fei-protocol/fei-protocol-core#16.

Description

commit allows anyone to commit purchased FGEN to a swap that will occur once the genesis group is launched. This commitment may be performed on behalf of other users, as long as the calling account has sufficient allowance:

code/contracts/genesis/GenesisGroup.sol:L87-L94

```
function commit(address from, address to, uint amount) external override onlyGenesisPeriod {
  burnFrom(from, amount);

committedFGEN[to] = amount;
  totalCommittedFGEN += amount;

emit Commit(from, to, amount);
}
```

The amount stored in the recipient's committedFGEN balance overwrites any previously-committed value. Additionally, this also allows anyone to commit an amount of "0" to any account, deleting their commitment entirely.

Recommendation

Ensure the committed amount is added to the existing commitment.

3.2 Purchasing and committing still possible after launch Critical

Resolution

This was addressed in fei-protocol/fei-protocol-core#11.

Description

Even after GenesisGroup.launch has successfully been executed, it is still possible to invoke GenesisGroup.purchase and GenesisGroup.commit.

Recommendation

Consider adding validation in GenesisGroup.purchase and GenesisGroup.commit to make sure that these functions cannot be called after the launch.

3.3 UniswapIncentive overflow on pre-transfer hooks Major

Resolution This was addressed in fei-protocol/fei-protocol-core#15.

Description

Before a token transfer is performed, Fei performs some combination of mint/burn operations via UniswapIncentive.incentivize:

code/contracts/token/UniswapIncentive.sol:L49-L65

```
function incentivize(
  address sender,
  address receiver,
  address operator,
  uint amountIn
) external override onlyFei {
    updateOracle();

  if (isPair(sender)) {
    incentivizeBuy(receiver, amountIn);
  }

  if (isPair(receiver)) {
      require(isSellAllowlisted(sender) || isSellAllowlisted(operator), "UniswapIncentive: Blocked F incentivizeSell(sender, amountIn);
  }
}
```

Both incentivizeBuy and incentivizeSell calculate buy/sell incentives using overflow-prone math, then mint / burn from the target according to the results. This may have unintended consequences, like allowing a caller to mint tokens before transferring them, or burn tokens from their recipient.

Examples []

incentivizeBuy calls getBuyIncentive to calculate the final minted value:

code/contracts/token/UniswapIncentive.sol:L173-L186

```
function incentivizeBuy(address target, uint amountIn) internal ifMinterSelf {
  if (isExemptAddress(target)) {
    return;
  }
    (uint incentive, uint32 weight,
    Decimal.D256 memory initialDeviation,
    Decimal.D256 memory finalDeviation) = getBuyIncentive(amountIn);
    updateTimeWeight(initialDeviation, finalDeviation, weight);
    if (incentive != 0) {
        fei().mint(target, incentive);
    }
}
```

getBuyIncentive calculates price deviations after casting amount to an int256, which may overflow:

code/contracts/token/UniswapIncentive.sol:L128-L134

```
function getBuyIncentive(uint amount) public view override returns(
    uint incentive,
    uint32 weight,
    Decimal.D256 memory initialDeviation,
    Decimal.D256 memory finalDeviation
) {
    (initialDeviation, finalDeviation) = getPriceDeviations(-1 * int256(amount));
```

Recommendation

Ensure casts in getBuyIncentive and getSellPenalty do not overflow.

3.4 BondingCurve allows users to acquire FEI before launch

Medium

Resolution

This was addressed in fei-protocol/fei-protocol-core#59

Description

BondingCurve.allocate allocates the protocol's held PCV, then calls _incentivize, which rewards the caller with FEI if a certain amount of time has passed:

code-update/contracts/bondingcurve/BondingCurve.sol:L180-L186

```
/// @notice if window has passed, reward caller and reset window
function _incentivize() internal virtual {
   if (isTimeEnded()) {
        _initTimed(); // reset window
        fei().mint(msg.sender, incentiveAmount);
   }
}
```

allocate can be called before genesis launch, as long as the contract holds some nonzero PCV. By force-sending the contract 1 wei, anyone can bypass the majority of checks and actions in allocate, and mint themselves FEI each time the timer expires.

Recommendation

Prevent allocate from being called before genesis launch.

3.5 Timed.isTimeEnded returns true if the timer has not been initialized Medium

Resolution This was addressed in fei-protocol/fei-protocol-core#62

Description

Timed initialization is a 2-step process:

- Timed.duration is set in the constructor: https://github.com/ConsenSys/fei-protocol-audit-2021-01/blob/d31114d834e62b4f3d4fa7b1c0b0c70fbff623a4/code-update/contracts/utils/Timed.sol#L15-L20
- Timed.startTime is set when the method _initTimed is called: https://github.com/ConsenSys/fei-protocol-audit-2021-01/blob/d31114d834e62b4f3d4fa7b1c0b0c70fbff623a4/code-update/contracts/utils/Timed.sol#L43-L46

Before this second method is called, isTimeEnded() calculates remaining time using a startTime of 0, resulting in the method returning true for most values, even though the timer has not technically been started.

Recommendation

If Timed has not been initialized, isTimeEnded() should return false, or revert

3.6 Overflow/underflow protection Medium

Resolution

This was partially addressed in fei-protocol/fei-protocol-core#17 by using SafeMath for the specific example given in the description.

Description

Having overflow/underflow vulnerabilities is very common for smart contracts. It is usually mitigated by using SafeMath or using solidity version 0.8 (after solidity 0.8 arithmetical operations already have default overflow/underflow protection).

In this code, many arithmetical operations are used without the 'safe' version. The reasoning behind it is that all the values are derived from the actual ETH values, so they can't overflow.

On the other hand, some operations can't be checked for overflow/underflow without going much deeper into the codebase that is out of scope:

code/contracts/genesis/GenesisGroup.sol:L131

```
uint totalGenesisTribe = tribeBalance() - totalCommittedTribe;
```

Recommendation

In our opinion, it is still safer to have these operations in a safe mode. So we recommend using SafeMath or solidity version ^0.8 compiler.

3.7 Unchecked return value for IWETH. transfer call Medium

Resolution

This was addressed in fei-protocol/fei-protocol-core#12.

Description

In EthUniswapPCVController, there is a call to IWETH.transfer that does not check the return value:

code/contracts/pcv/EthUniswapPCVController.sol:L122

```
weth.transfer(address(pair), amount);
```

It is usually good to add a require-statement that checks the return value or to use something like safeTransfer; unless one is sure the given token reverts in case of a failure.

Recommendation

Consider adding a require-statement or using safeTransfer.

3.8 GenesisGroup.emergencyExit remains functional after launch Medium

Resolution

This was partially addressed in fei-protocol/fei-protocol-core#14 and fei-protocol/fei-protocol-core#13 by addressing the last two recommendations.

Description

emergencyExit is intended as an escape mechanism for users in the event the genesis launch method fails or is frozen. emergencyExit becomes callable 3 days after launch is callable. These two methods are intended to be mutually-exclusive, but are not: either method remains callable after a successful call to the other.

This may result in accounting edge cases. In particular, emergencyExit fails to decrease totalCommittedFGEN by the exiting user's commitment:

code/contracts/genesis/GenesisGroup.sol:L185-L188

```
burnFrom(from, amountFGEN);
committedFGEN[from] = 0;

payable(to).transfer(total);
```

As a result, calling launch after a user performs an exit will incorrectly calculate the amount of FEI to swap:

code/contracts/genesis/GenesisGroup.sol:L165-L168

```
uint amountFei = feiBalance() * totalCommittedFGEN / (totalSupply() + totalCommittedFGEN);
if (amountFei != 0) {
  totalCommittedTribe = ido.swapFei(amountFei);
}
```

Recommendation

- Ensure launch cannot be called if emergencyExit has been called
- Ensure emergencyExit cannot be called if launch has been called
- ullet In ${\tt emergencyExit}$, reduce ${\tt totalCommittedFGEN}$ by the exiting user's committed amount

3.9 Unchecked return value for transferFrom calls Medium

Resolution

This was addressed in fei-protocol/fei-protocol-core#12.

Description

There are two transferFrom calls that do not check the return value (some tokens signal failure by returning false):

code/contracts/pool/Pool.sol:L121

```
stakedToken.transferFrom(from, address(this), amount);
```

code/contracts/genesis/IDO.sol:L58

```
fei().transferFrom(msg.sender, address(pair), amountFei);
```

It is usually good to add a require-statement that checks the return value or to use something like safeTransferFrom; unless one is sure the given token reverts in case of a failure.

Recommendation

Consider adding a require-statement or using safeTransferFrom.

3.10 GovernorAlpha proposals may be canceled by the proposer, even after they have been accepted and queued Minor

Resolution

This was addressed in fei-protocol/fei-protocol-core#61

Description

GovernorAlpha allows proposals to be canceled via cancel. To cancel a proposal, two conditions must be met by the proposer:

- The proposal should not already have been executed: https://github.com/ConsenSys/fei-protocol-audit-2021-01/blob/d31114d834e62b4f3d4fa7b1c0b0c70fbff623a4/code-update/contracts/dao/GovernorAlpha.sol#L206-L208
- The proposer must have under proposalThreshold() TRIBE balance:
 https://github.com/ConsenSys/fei-protocol-audit-202101/blob/d31114d834e62b4f3d4fa7b1c0b0c70fbff623a4/codeupdate/contracts/dao/GovernorAlpha.sol#L210-L211

The latter condition is completely under the control of the proposer, meaning that a proposer may cancel proposals in any of these states: Pending, Active, Canceled, Defeated, Succeeded, Queued, Expired.

Recommendation

Prevent proposals from being canceled unless they are in the Pending or Active states.

3.11 Pool: claiming to the pool itself causes accounting issues Minor

Resolution

This was addressed in fei-protocol/fei-protocol-core#57

Description

In Pool.sol, claim(address from, address to) is used to claim staking rewards and send them to a destination address to:

code-update/contracts/pool/Pool.sol:L229-L238

```
function _claim(address from, address to) internal returns (uint256) {
    (uint256 amountReward, uint256 amountPool) = redeemableReward(from);
    require(amountPool != 0, "Pool: User has no redeemable pool tokens");

    _burnFrom(from, amountPool);
    _incrementClaimed(amountReward);

    rewardToken.transfer(to, amountReward);
    return amountReward;
}
```

If the destination address to is the pool itself, the pool will burn tokens and increment the amount of tokens claimed, then transfer the reward tokens to itself.

Recommendation

Prevent claims from specifying the pool as a destination.

3.12 Assertions that can fail Winor

Description

In UniswapSingleEthRouter there are two assert-statements that may fail:

code/contracts/router/UniswapSingleEthRouter.sol:L21

```
assert(msg.sender == address(WETH)); // only accept ETH via fallback from the WETH contract
```

code/contracts/router/UniswapSingleEthRouter.sol:L48

```
assert(IWETH(WETH).transfer(address(PAIR), amountIn));
```

Since they do some sort of input validation it might be good to replace them with require-

statements. I would only use asserts for checks that should never fail and failure would constitute a bug in the code.

Recommendation

Consider replacing the assert-statements with require-statements. An additional benefit is that this will not result in consuming all the gas in case of a violation.

3.13 Simplify API of GenesisGroup. purchase Minor

Description

The API of GenesisGroup.purchase could be simplified by not including the value parameter that is required to be equivalent to msg.value:

code/contracts/genesis/GenesisGroup.sol:L79

```
require(msg.value == value, "GenesisGroup: value mismatch");
```

Using msg.value might make the API more explicit and avoid requiring msg.value == value. It can also save some gas due to fewer inputs and fewer checks.

Recommendation

Consider dropping the value parameter and changing the code to use msg.value instead.

4 Infrastructure Security Assessment

Each issue has an assigned severity:

- Minor issues are subjective in nature. They are typically suggestions around best practices or readability. Code maintainers should use their own judgment as to whether to address such issues.
- Medium issues are objective in nature but are not security vulnerabilities. These should be addressed unless there is a clear reason not to.
- Major issues are security vulnerabilities that may not be directly exploitable or may require certain conditions in order to be exploited. All major issues should be addressed.
- Critical issues are directly exploitable security vulnerabilities that need to be fixed.

4.1 Clickjacking and Missing Content Security Policy Major Fixed

Resolution

After multiple iterations, the following Content Security Policy has been put into effect:

```
default-src 'self';
script-src 'self' 'unsafe-inline' https://www.googletagmanager.com;
object-src 'none';
style-src 'self' 'unsafe-inline';
img-src 'self';
media-src 'none';
frame-src 'none';
font-src 'self';
connect-src 'self'
  https://api.amplitude.com
  https://eth-ropsten.alchemyapi.io
  https://eth-mainnet.alchemyapi.io
  https://api.thegraph.com;
frame-ancestors 'none'
```

The CSP is transmitted through the following headers:

- Content-Security-Policy
- X-Content-Security-Policy
- X-WebKit-CSP

as well as through corresponding meta HTML tags. Additionally, the following frame-busting JavaScript code has been added to prevent Clickjacking attacks in the unlikely event that existing CSP measures fail or are bypassed:

```
<script>
if (self == top) {
  document.documentElement.style.display = 'block ';
} else {
  top.location = self.location;
}
</script>
```

Description

A content security policy (CSP) provides an added layer of protection against cross-site scripting (XSS), clickjacking, and other client-side attacks that rely on executing malicious content in the context of the website.

Specifically, the lack of a content security policy allows an adversary to perform a clickjacking attack by including the target URL (such as app.fei.money) in an iframe element on their site. The attacker then uses one or more transparent layers on top of the embedded site to trick a user into performing a click action on a different element.

This technique can be used to spawn malicious Metamask dialogues, tricking users into thinking that they are signing a legitimate transaction.

Affected Assets

All S3-hosted web sites.

Recommendation

It is recommended to add content security policy headers to the served responses to prevent browsers from embedding Fei-owned sites into malicious parent sites. Furthermore, CSP can be used to limit the permissions of JavaScript and CSS on the page, which can be used to further harden the deployment against a potential compromise of script dependencies.

It should be noted that security headers should not only be served from Cloudfront but any public-facing endpoint. Otherwise, it will be trivial for an attacker to circumvent the security headers added by Cloudfront, e.g. by embedding the index.html file directly from the public-facing S3 bucket URL.

Besides CSP headers, clickjacking can also be mitigated by directly including frame-busting JavaScript code into the served page.

4.2 S3 Buckets Cleartext Communication Medium Fixed

Resolution

Direct access to S3 buckets through s3.amazonaws.com is now rejected, while unencrypted HTTP traffic to the previously affected assets now consistently redirects to the HTTPS equivalents.

Description

The system's S3 buckets are configured to allow unencrypted traffic:

```
$ curl -v http://fei.money.s3.amazonaws.com/index.html
  Trying 52.219.112.162:80...
* TCP_NODELAY set
* Connected to fei.money.s3.amazonaws.com (52.219.112.162) port 80 (https://github.com/ConsenSys/fei-protocol-audit-20
> GET /index.html HTTP/1.1
> Host: fei.money.s3.amazonaws.com
> User-Agent: curl/7.68.0
> Accept: */*
* Mark bundle as not supporting multiuse
< HTTP/1.1 200 OK
< x-amz-id-2: 0QtzqEhGn7gHUjjiAxpni0MXKQ101ouT6Tp8iQG2EfvlKbg0ZgEbDdkQrJrJL20yJF1VyZkPjjU=</pre>
< x-amz-request-id: D6250FE8F76E84F0
< Date: Tue, 09 Feb 2021 13:07:54 GMT
< Last-Modified: Mon, 11 Jan 2021 20:38:09 GMT
< ETag: "ec826fa83693f3db3a989fcbeb5adef1"
< Accept-Ranges: bytes
< Content-Type: text/html
< Content-Length: 3675
< Server: AmazonS3
< ...
```

Affected Assets

```
arn:aws:s3:::ropsten-app.fei.money/*
arn:aws:s3:::www.fei.money/*
arn:aws:s3:::feiprotocol.com/*
arn:aws:s3:::www.app.fei.money/*
arn:aws:s3:::www.ropsten-app.fei.money/*
arn:aws:s3:::app.fei.money/*
```

arn:aws:s3:::fei.money/*

Recommendation

It is recommended to enforce encryption of data in transit using TLS certificates. To accomplish this, the aws:SecureTransport can be set in the S3 bucket's policies.

4.3 Missing Log Aggregation Medium Fixed

Resolution

CloudFrond and CloudTrail have been enabled. These components send endpoint-related and organizational log messages into S3 buckets where they can be queried using AWS Athena. The security review process section of this report contains sample queries for Athena.

Description

There is no centralized system that gathers operational events of AWS stack components. This includes S3 server access logs, configuration changes, as well as Cloudfront-related logging.

Recommendation

It is recommended to enable CloudTrail for internal log aggregation as it integrates seamlessly with S3, Cloudfront, and IAM. Furthermore, regular reviews should be set up where system activity is checked to detect suspicious activity as soon as possible.

4.4 Enforce Strict Transport Security Medium Fixed

Resolution

All domains in scope now ship with the following header:

strict-transport-security: max-age=63072000; includeSubdomains

Description

The HTTP Strict-Transport-Security response header (often abbreviated as HSTS) lets a web site tell browsers that it should only be accessed using HTTPS, instead of using HTTP. This prevents attackers from stripping TLS certificates from connections and removing encryption.

Recommendation

It is recommended to deliver all responses with the strict-Transport-Security header. In an S3-Cloudfront setup, this can be achieved using Lambda@Edge lambda functions.

4.5 Server Information Leak Minor

Description

Responses from the fei.money domain and related assets leak server information in their response headers. This information can be used by an adversary to prepare more sophisticated attacks tailored to the deployed infrastructure.

Note: At the time of reporting, this issue was deemed not possible to fix due to technical limitations on AWS-hosted static sites using S3 and CloudFront.

Examples

```
$ curl -I https://fei.money/static/media/
HTTP/2 404
x-amz-error-code: NoSuchKey
x-amz-error-message: The specified key does not exist.
x-amz-error-detail-key: static/media/index.html
date: Tue, 09 Feb 2021 13:49:34 GMT
server: AmazonS3
x-cache: Error from cloudfront
via: 1.1 fa133af2508a341e1ff6bfff526ba095.cloudfront.net (CloudFront)
x-amz-cf-pop: TXL52-C1
x-amz-cf-id: x0eNuDCrilaFg0T3fz4g1CpdRIfFCxBta7Pif4wexsXpN3weVLv7uw==
```

Recommendation

It is recommended to remove any headers that hint at server technologies and are not directly required by the frontend.

4.6 Missing Route53 Domain Lock Minor Fixed

```
Resolution

A transfer lock on both the feiprotocol.com and fei.money domains has been requested.
```

Description

Domain registrars often give customers the option to lock a domain. This prevents unauthorized parties from transferring it to another registrar, either through malicious interaction with the registrar itself, or compromised domain owner credentials. No domain currently has a lock enabled.

Affected Assets

- fei. money
- feiprotocol.com

Recommendation

It is recommended to set a lock for the affected domains, assuming that the registrar allows domain locks:

- 1. Sign in to the AWS Management Console and open the Route 53 console at https://console.aws.amazon.com/route53/.
- 2. In the navigation pane, choose Registered Domains.

- 3. Choose the name of the domain that you want to update.
- 4. Choose Enable (to lock the domain) or Disable (to unlock the domain).
- Choose Save.

4.7 Weak IAM Password Policy Minor Fixed

Resolution

This has been fixed by the client with the following notes:

- Enforced 14 character password length
- Enabled 90 day password expiration
- Prevent password reuse
- Require one uppercase, one lowercase, one number, one non-alphanumeric character
- Require 2FA on all users via this doc and this post (Create new Force_MFA policy, attach it to the new Engineers group, and then assign all users (including Dominik) to this group
- Also requiring 2FA on command line access. Using src/infra/aws-token.sh for generating
 the credentials and putting them in ~/.aws/config

Description

The password policy for IAM users currently does not enforce the use of strong passwords, multi-factor authentication, and regular password rotation.

Currently, only a minimum password length of 8 is enforced.

Recommendation

- Require a minimum password length of 14
- Set a password expiration policy of at most 90 days
- Disallow the reuse of passwords
- Enable mandatory multi-factor authentication with a virtual app

4.8 Review Access Key Expiration Minor Fixed

Resolution

This issue is considered resolved with the implementation of regular security review meetings.

Description

It is recommended to only create access keys when absolutely necessary. There should be no

access keys given out to root users. Instead, temporary security credentials (IAM Roles) should be created.

Recommendation

It is recommended to read the Best practices for managing AWS access keys and incorporate the security practices where reasonable.

4.9 Dependency Security Minor Fixed

Resolution

This issue has been resolved by implementing Snyk for continuous dependency security scanning. This allows the developers to review potential risks of included packages and receiving automated pull requests with fixes if necessary.

Furthermore, a manual review of select dependencies has been conducted by the penetration tester without significant, actionable results. The following dependencies have been checked:

- bignumber
- numeral
- validator
- web-vitals

Description

The Yarn audit feature currently finds two low-severity dependency issues:

- Prototype pollution in ini a dependency of react-scripts
- Insecure Credential Storage in web3

yarn audit v1.22.4		
low	Prototype Pollution	1
Package	ini	1
Patched in	>1.3.6	I
Dependency of	react-scripts	I
	react-scripts > react-dev-utils > global-modules > global-prefix > ini	
More info	https://www.npmjs.com/advisories/1589	I
low	Insecure Credential Storage	1
Package	web3	I
Patched in	No patch available	I
Dependency of	web3	I
Path	web3	I
More info	https://www.npmjs.com/advisories/877	I
2 vulnerabilities found - Packages audited: 1963 Severity: 2 Low Done in 11.90s.		
1		Þ

Recommendations

It is recommended to apply the ini patch, which is already available. For web3, it is recommended to monitor the repository's Github issue https://github.com/ConsenSys/fei-protocol-audit-2021-01/issues/2739 and upgrade as soon as a fix is available.

For additional dependency security, it is recommended to integrate a security monitoring service. Snyk has a free plan which allows unlimited tests on public repositories, and 200 tests per month for private ones. A bot will automatically add a pull request to bump vulnerable dependency versions.

It should be noted that the quality and reliability of such automated contributions are highly dependent on the quality of the test suite. It is recommended to build strict tests around core functionality and expected dependency behaviour to detect breaking changes as soon as possible.

5 Security Review Process

In an additional effort to achieve security-in-depth, it is recommended to implement a schedule or recurring security review meetings. The goal of these meetings is to complete a checklist to enforce security best-practices, as well as find anomalies in the system as soon as possible to commence mitigation and investigations.

This section outlines recommendations for the contents of such a checklist. It should be noted that security requirements are likely to change, and thus, this list should be treated as a working document as the project's infrastructure and attack surface change.

5.1 CloudTrail Anomalies

Event filter query template:

Event Filter	Description
eventname = 'ConsoleLogin'	Sign-in activity
eventname = 'AddUserToGroup'	User added to group
eventname = 'ChangePassword'	User password change
eventname LIKE '%AccessKey%'	Key management events
eventname LIKE '%MFADevice'	MFA deactivation/deletion/resync
<pre>eventname = 'StopLogging'</pre>	Logging stopped
eventname LIKE '%BucketPolicy%'	Bucket policy activity
eventname LIKE '%GroupPolicy%'	Group policy activity
eventname LIKE '%UserPolicy%'	User policy activity
eventname LIKE '%RolePolicy%'	Role policy activity

Aggregate statistics about failed authentication and user authorization attempts can be gathered with the following query:

For investigative purposes or the goal of covering new infrastructure components, it might be necessary to add more event names to the review process. AWS does not provide a comprehensive list of event names per stack component. An external list of CloudTrail event names is available on the GorillaStack blog.

Note: In case of issues with Athena or ingestion into the database, CloudTrail allows users to view the unfiltered event history for a user-specified time range as well. Particularly notable is the ability to filter by resource types, of which the following are relevant to the Fei AWS infrastructure:

AWS::S3::Bucket
AWS::CloudTrail::Trail
AWS::IAM::AccessKey
AWS::IAM::MfaDevice
AWS::IAM::Group
AWS::IAM::Policy
AWS::IAM::Role

5.2 Cloudfront Endpoint Anomalies

Top 10 endpoints hit in a given time frame:

This query can be filtered further by adding AND status = 500 or a similar condition to find suspicious response codes.

5.3 Route53 CNAME Review

Subdomain takeover vulnerabilities occur when a subdomain is pointing to a service, e.g. a previously deleted CloudFront endpoint or S3 bucket. This allows an attacker to set up a page on the service that was being used and point their page to that subdomain. Especially with wildcard certificates on the system, e.g. *.fei.money*, this can lead to an exploitation of user trust and enables attacks that can result in reputational and financial loss.

It is recommended that DNS records in Route53 are reviewed regularly and removed as soon as the underlying resource is decommissioned.

5.4 External Monitoring and Notifications

Beyond manual checks, it is recommended that a service such as Assertible is used. This will allow the development team to detect unavailable endpoints and enforce regularly-checked assertions, such as proper return codes or page content. Furthermore, such a service should integrate other means of communication such as Slack notifications, SMS messages, or arbitrary webhook calls to notify an on-duty developer as quickly as possible.

Appendix 1 - Disclosure

ConsenSys Diligence ("CD") typically receives compensation from one or more clients (the "Clients") for performing the analysis contained in these reports (the "Reports"). The Reports may be distributed through other means, including via ConsenSys publications and other distributions.

The Reports are not an endorsement or indictment of any particular project or team, and the Reports do not guarantee the security of any particular project. This Report does not consider, and should not be interpreted as considering or having any bearing on, the potential economics of a token, token sale or any other product, service or other asset. Cryptographic tokens are emergent technologies and carry with them high levels of technical risk and uncertainty. No Report provides any warranty or representation to any Third-Party in any respect, including regarding the bugfree nature of code, the business model or proprietors of any such business model, and the legal compliance of any such business. No third party should rely on the Reports in any way, including for the purpose of making any decisions to buy or sell any token, product, service or other asset. Specifically, for the avoidance of doubt, this Report does not constitute investment advice, is not intended to be relied upon as investment advice, is not an endorsement of this project or team, and it is not a guarantee as to the absolute security of the project. CD owes no duty to any Third-Party by virtue of publishing these Reports.

PURPOSE OF REPORTS The Reports and the analysis described therein are created solely for Clients and published with their consent. The scope of our review is limited to a review of Solidity code and only the Solidity code we note as being within the scope of our review within this report. The Solidity language itself remains under development and is subject to unknown risks and flaws. The review does not extend to the compiler layer, or any other areas beyond Solidity that could present security risks. Cryptographic tokens are emergent technologies and carry with them high levels of technical risk and uncertainty.

CD makes the Reports available to parties other than the Clients (i.e., "third parties") - on its website. CD hopes that by making these analyses publicly available, it can help the blockchain ecosystem develop technical best practices in this rapidly evolving area of innovation.

LINKS TO OTHER WEB SITES FROM THIS WEB SITE You may, through hypertext or other computer links, gain access to web sites operated by persons other than ConsenSys and CD. Such hyperlinks are provided for your reference and convenience only, and are the exclusive responsibility of such web sites' owners. You agree that ConsenSys and CD are not responsible for the content or operation of such Web sites, and that ConsenSys and CD shall have no liability to you or any other person or entity for the use of third party Web sites. Except as described below, a hyperlink from this web Site to another web site does not imply or mean that ConsenSys and CD endorses the content on that Web site or the operator or operations of that site. You are solely responsible for determining the extent to which you may use any content at any other web sites to which you link from the Reports. ConsenSys and CD assumes no responsibility for the use of third party software on the Web Site and shall have no liability whatsoever to any person or entity for the accuracy or completeness of any outcome generated by such software.

TIMELINESS OF CONTENT The content contained in the Reports is current as of the date appearing on the Report and is subject to change without notice. Unless indicated otherwise, by ConsenSys and CD.

Request a Security Review Today

Get in touch with our team to request a quote for a smart contract audit.

CONTACT US



AUDITS FUZZING SCRIBBLE BLOG TOOLS RESEARCH ABOUT CONTACT CAREERS PRIVACY POLICY

Subscribe to Our Newsletter

Stay up-to-date on our latest offerings, tools, and the world of blockchain security.

